



(a) Strip convolution, adopted from HPLNet



Factorised kernel: 6 instead of 9 weights per channel pair, strong response to horizontal and vertical line segments

(b) Channel shift, adopted from HPLNet



25% of channels in four groups, each moved by 2 pixels; no weights and no multiply operations

(c) Encoder and decoder for a 256×256 input

