




Measured on Intel(R) Core(TM) i5-1035G4 CPU @ 1.10GHz (4 threads), PyTorch 2.13.0, batch 1, median of 3 to 7 runs. Reference models untrained: compute cost only.
† Full 1024^2 pass exceeded this machine's free RAM: FLOPs counted at 256^2 and scaled $\times 16$ (exact for fully convolutional nets); latency timed as four 512^2 crops.